

Q & A related to Assignment 1

1. The `pstat` command requires us to list the information of `utime` and `stime` of a process. I tried to find this information in `/proc/pid/stat` file, but found `utime` and `stime` are always zero. Why this happened?

Answer: The unit of `utime/stime` is in `CLOCK TICKS`.

In Linux system, `sysconf(_SC_CLK_TCK)` returns 100 by default (although the default value might change since modern computers may dynamically change the value depending on current work load). So OS keeps time measurements in units of 100 cycles per second, or each system tick every 10 ms. In other words, if the value of `utime` is 4, then the value translate to $4/100 = 0.4$ seconds, which is 40 ms. 0.4 seconds (or 40 ms) is the value that your code should print out.

For the above reason, you should cast the return value as float instead of int. In the above example, `int 4 / int 100` will return 0, which is wrong.

Then why are `utime` and `stime` zero nearly always? This phenomenon is because your background process is nearly always sleeping. If you use `usleep()` rather than using `sleep()` to let linux schedule your process more frequently (like `usleep(100)`, which let linux to schedule your process every 100 micro seconds), then you should be able to see that `utime` and `stime` are not zero anymore.

2. The `pstat` command requires us to list the information of `RSS` of a process. But both `/proc/pid/stat` and `/proc/pid/status` have the information of `RSS`, which one should I use?

Answer: You should use the one in `/proc/pid/stat`, the one listed in `/proc/pid/status` is `VMRSS`, which means number of pages the process has in virtual memory. Since virtual memory = the part in physical memory + the part on disk, thus `VMRSS` should be greater than `RSS`.

3. Where can I find the information (state, `utime` etc.) that `pstat` command required?

Answer: All the information `pstat` required can be found in `/proc/pid/stat` and `/proc/pid/status` two pseudo files. Their specifications have been posted in `connex` → `resource` → Assignment 1

4. Should PMan care about the background processes that were not executed by PMan `bg`?

Answer: No, you do not need to handle those background processes that were not executed by PMan `bg`.

5. Should PMan indicate to the user when a background process was terminated by a `kill` command outside of PMan?

Answer: Yes. A background process can be killed either inside PMan by `bgkill` or outside PMan e.g., by using command line `% kill pid` from another terminal. Either case, PMan should tell the user that the background process has been terminated. Please be reminded again that your code does not need to care about those background processes that were not executed by PMan `bg`.

To make your life easy, we do not require that your PMan immediately reports to the user when a background terminates. Such events can be reported at a later time when a user types `bglist` in PMan. Of course, you can also implement PMan in a way that it immediately reports the termination of a background process, but this is not mandatory.

6. How to implement the requirement stated in Question 5?

Answer: use system call `waitpid` with `-1` as the first parameter, e.g.,

```
pid = waitpid(-1, &p_status, WNOHANG)
```

where `-1` means wait for any child process, `WNOHANG` means "return immediately if no child has exited."

Note that `waitpid` returns the process ID of the terminated process whose status was reported. If unsuccessful, a `-1` is returned.

You can further check the value in `p_status` to tell if a child process was killed or has exited. E.g.,

```
if (WIFSIGNALED(p_status)) {
    printf("Process %d was killed\n", pid);
    remove_process_from_list(pid);
}
if (WIFEXITED(p_status)) {
    printf("Process %d exits\n", pid);
    remove_process_from_list(pid);
}
```

7. After I type `bg test` in PMan where `test` is a fake command, the child process calls `execvp` to execute `test` and the call fails. Nevertheless, my parent process adds and lists `test` as a background process even if it fails. How can I solve this problem?

Ans: First of all, your child process needs to capture the returned value of `execvp`. If the call fails, the child process needs some way to notify this failure to the parent process if you maintain the list of background processes in the parent process. In other words, you need some communication mechanism between the child and the parent processes. There are different ways that you can use for this purpose:

(a) Option 1: In your child process, when *execvp* fails, the child process exit with an error code, e.g.,

```
if (execvp(input[0], input) < 0) {  
    printf("Error: execution of %s failed \n", input[0]);  
    exit(-1);  
}
```

In the parent process, before you add the background process to the list, you can use *waitpid* with *WNOHANG* to check the error code from the child. Since *WNOHANG* is used to avoid blocking the parent process, there is a chance that *waitpid* might miss the error code and thus the background process (with the process id being the child process id, and the name of the background process being *test*) is added to the list. (This could happen if the child process runs *exit(-1)* after the parent process calls *waitpid*.) Nevertheless, when you type *bglis*t in PMan afterwards, the background process (*test*) must be shown as terminated, because it actually never got executed.

Of course, this is not a perfect solution, because *bglis*t will show *test* as a terminated background process even if it never got executed. However, you will not lose mark for such minor inconsistency.

(b) Option 2: The list of background process is stored in a shared memory space (refer to the IPC demo code). In this way, the child and parent process have an easy way to coordinate. In this case, inserting new background process into the background process list is done at the child process.