

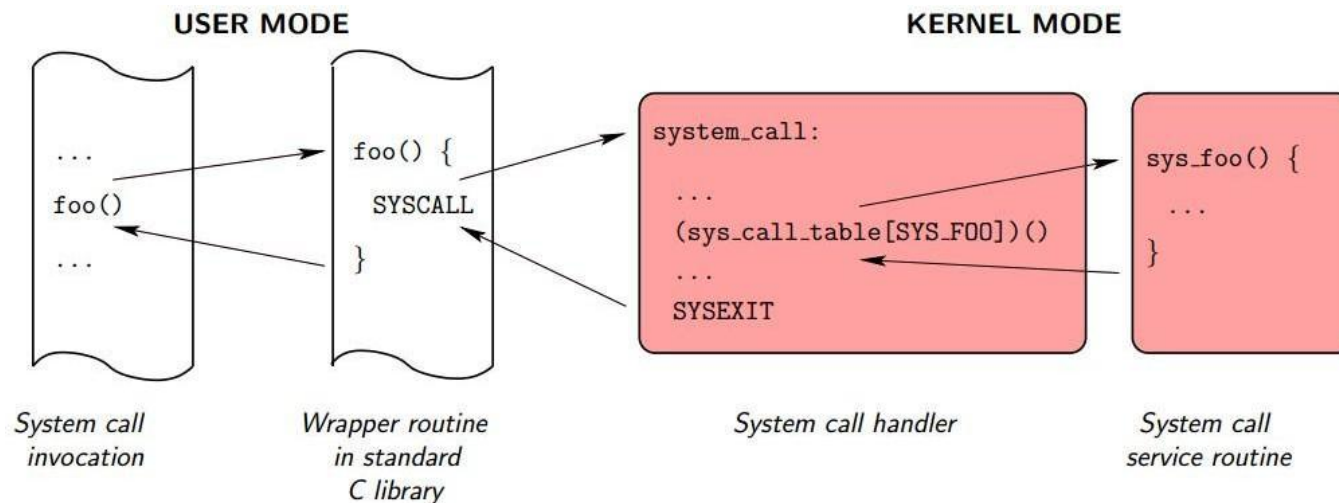
OS Tutorial 2: System Calls

Huan Wang

huanwang@uvic.ca

System Calls (1)

- * **kernel interfaces:** services provided by the OS kernel
- * Use system calls like library functions, including necessary **header files** (e.g., **#include <unistd.h>**)



System Calls (2) – Process Mgmt.

- * **fork()**: create a new (child) process
- * **exec*()**: a family of functions: execute other programs from a process
- * **kill()**: send signals to process
- * **wait()/waitpid()**: wait for a (child) process to change state
- * Use **man** to check the details

fork()

- * By calling fork(), a process can spawn a child process.
- * The child **duplicates** the parent process, so it is **almost** identical to the parent (except **PID**, etc.)
- * After fork(), both the parent and child execute the same program
- * Prototype: `pid_t pid = fork();` (**pid_t** is an alias to an integer type)
- * Header files: `#include <unistd.h>` & `#include<sys/types.h>`
- * Return values:
 - * `pid > 0`: succeed, in parent process
 - * `pid == 0`: succeed, in child process
 - * `pid < 0`: fail.
- * For more details: **\$man 2 fork**

exec*() (1)

- * A family of six functions:
 - * `int execl(char *path, char *arg, ...);`
 - * `int execv(char *path, char *argv[]);`
 - * `int execlp(char *path, char *arg, ..., char *envp[]);`
 - * `int execve(char *path, char *argv[], char *envp[]);`
 - * `int execlp(char *file, char *arg, ...);`
 - * `int execvp(char *file, char *argv[]);`
- * What *l*, *v*, *e*, and *p* mean:
 - * *l* means an argument **list**,
 - * *v* means an argument **vector**,
 - * *e* means an **environment** vector, and
 - * *p* means a environment **path**.
- * For more details: **\$man 3 exec**

exec*() (2)

- * load and run a new program so as to replace the current process
- * Upon success, **exec()** **never** returns to the caller unless the call **failed**.
 - * Failed reasons: non-existent file (bad path) or bad permissions

kill()

- * Send signals to a process specified by PID
- * Prototype:
 - * `int kill(pid_t pid, int signal);`
 - * E.g., `int retVal = kill(child_pid, SIGTERM);`
- * Header files: `#include <sys/types.h>` & `#include <signal.h>`
- * Return values:
 - * On success: `retVal = 0`
 - * On error: `retVal = -1`
- * Linux signals:
 - * **SIGSTOP, SIGCONT, SIGTERM, SIGKILL**, etc.
 - * **\$ man 7 signal**

wait

- * Forces the parent to suspend execution and wait for its children or a specific child to change states/state.

Two forms:

- * `wait()`: parent waits for any child process

- * Prototype:

- * `pid_t wait(int *status);`

- * Header file: `#include <sys/types.h> & #include <sys/wait.h>`

- * Return values:

- * On success: PID of the exited process
 - * On error: -1

waitpid()

- * Parent waits for a state change of a child with given PID
- * `pid_t waitpid(pid_t pid, int *wstatus, int options);`
- * Options:
 - * If 0, then waits until the specified child return
 - * **WNOHANG** - return immediately if no child has exited
 - * **WUNTRACED** - also returns if a child has stopped
 - * **WCONTINUED** - also returns if a stopped child has been resumed by SIGCONT
- * return value:
 - * return the PID of the child whose state has changed
 - * return 0 if WNOHANG was specified but child have not yet changed state
 - * return -1 on **error**

Note: the parent can only wait its **direct child** (**NOT** grandchild) * For more details: **\$man 2 wait**

exit()

- * **Gracefully** terminates process execution: clean up and release resources; puts the process into **zombie** state.
- * Prototype: `void exit(int status);`
- * Header file: `#include <stdlib.h>`
- * **exit()** specifies a return value from the process (i.e., status of the dead process), which a parent process may need to examine.
- * **_exit()** call is another possibility of quick death without cleanup.
- * For more details: `$ man 3 exit` & `$ man 2 _exit`

Background Execution

- * Switch a program from the foreground to the background or vice-versa.
- * Linux commands:
 - * **Ctrl + C**: Terminate the foreground process and return to Shell
 - * **Ctrl + Z**: Suspend the foreground process, send it to background and return to Shell
 - * **&**: Let the program run at background
 - * **fg [num]**: Move the process with job ID=num to foreground
 - * **bg [num]**: Move the process with job ID=num to background
- * Online tutorial:
<http://www.thegeekstuff.com/2010/05/unix-background-job/>

Outline

- * **System call (Questions?)**

Contributors:

- * Cheng Chen, Guoming Tang, Yongjun Xu
- * Huan Wang, Changli Zhang